

Here's a beginner-friendly overview and code examples for three core Machine Learning algorithms:

Machine Learning Algorithms

◆ 1. K-Nearest Neighbors (KNN) — *Supervised Learning*

What it does:

- Classifies new data points based on the majority class of its k nearest neighbors in the training data.
- Works well for classification and regression (commonly used for classification).

How it works:

- Calculate distance (e.g., Euclidean) from the test point to all training points.
- Pick the k closest points.
- Assign the class that is most common among these k.

Example (using sklearn):

```
from sklearn.neighbors import KNeighborsClassifier
```

```
# Sample data
```

```
X = [[1, 2], [2, 3], [3, 3], [6, 5], [7, 7]]
```

```
y = ['A', 'A', 'A', 'B', 'B']
```

```
# Train the model
```

```
model = KNeighborsClassifier(n_neighbors=3)
```

```
model.fit(X, y)
```

```
# Predict new value
```

```
print("Predicted Class:", model.predict([[4, 4]])) # Should classify between A and B
```

◆ 2. K-Means Clustering — *Unsupervised Learning*

What it does:

- Groups data into k clusters where each point belongs to the cluster with the nearest mean (centroid).

- Unlabeled data.

 **How it works:**

- Randomly select k centroids.
- Assign points to the nearest centroid.
- Recompute centroids.
- Repeat until convergence.

 **Example:**

```
from sklearn.cluster import KMeans
```

```
import numpy as np
```

```
# Sample data
```

```
X = np.array([[1, 2], [1, 4], [1, 0],  
             [10, 2], [10, 4], [10, 0]])
```

```
# Apply KMeans
```

```
kmeans = KMeans(n_clusters=2, random_state=0)  
kmeans.fit(X)
```

```
print("Cluster Centers:\n", kmeans.cluster_centers_)  
print("Labels:", kmeans.labels_)
```

◆ **3. Linear Regression — *Supervised Learning***

 **What it does:**

- Models the relationship between independent variable(s) X and a dependent variable y using a straight line.
- Predicts continuous values.

 **How it works:**

- Finds the best-fit line:

$y=mx+cy = mx + c$

 **Example:**

```
from sklearn.linear_model import LinearRegression
```

```
import numpy as np

# Training data
X = np.array([[1], [2], [3], [4], [5]])
y = np.array([2, 4, 5, 4, 5])

# Train the model
model = LinearRegression()
model.fit(X, y)

# Predict
print("Predicted:", model.predict([[6]])) # Predict y when x=6
```

Summary Table

| Algorithm | Type | Use Case | Output Type |
|-------------------|--------------|-------------------------|-------------|
| KNN | Supervised | Classification | Class Label |
| K-Means | Unsupervised | Grouping / Segmentation | Cluster ID |
| Linear Regression | Supervised | Predicting Values | Continuous |

Would you like a small project that combines these algorithms using real-world data (e.g., Titanic, Iris, or customer segments)?